

working memory (KB)

- attention scores ($HL^2 \times 4 \text{ B}$)
- KV cache ($2Ld_{\text{model}} \times 4 \text{ B}$)
- ours: 6-layer recurrent state

16 MB

4 MB

1 MB

256 KB

64 KB

16 KB

32

64

128

256

512

1024

context length L (tokens)

crossover $L \approx 55$ tokens

48 KB, flat in L

evaluated window: 32 tokens
($W = 64$ cycles, patch 2)
attention 16 KB < our 48 KB state

the claim is the slope, not today's value

